

Project 1

Mike Keating, Hayden Morgan

Project 1

Setting Things Up

Creating the Repo

- GitHub repo created by Mike
- RStudio project created
- Hayden added as a collaborator and membership accepted
- Format set to PDF

Collaboration Workflow

- Task distribution and timeline established
- Decided to each work on own branches

.qmd Format

All messages and warnings that come from librarying packages should be turned off using the appropriate code chunk option.

```
library("tidyverse")  
library("ggplot2")
```

First Steps

Question 1: Selecting Columns

Read in one section of the data. This data is available at <https://www4.stat.ncsu.edu/~online/datasets/EDU01a.csv>.

Select only the following columns:

- Area_name (rename area_name)
- STCOU
- Any column that ends in “D”

#NOTE: EDU01a, EDU01b, divisions, and Mastdata files are all in the project folder.

```
edu01a <- read_csv("data/EDU01a.csv",  
                  col_select = c(Area_name, STCOU, ends_with("D")),  
                  show_col_types = FALSE) |>  
  rename(  
    area_name = Area_name  
  )
```

Display the first 5 rows of your new data set to show that you created this correctly. Note: Do not save over your new data set with just the first 5 rows, simply just show the first 5 rows.

```
head(edu01a, 5)
```

```
# A tibble: 5 x 12  
  area_name      STCOU EDU010187D EDU010188D EDU010189D EDU010190D EDU010191D  
  <chr>         <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>  
1 UNITED STATES 00000    40024299   39967624   40317775   40737600   41385442  
2 ALABAMA      01000     733735    728234     730048     728252     725541  
3 Autauga, AL   01001      6829      6900       6920       6847       7008  
4 Baldwin, AL  01003     16417     16465     16799     17054     17479  
5 Barbour, AL  01005      5071      5098      5068      5156      5173  
# i 5 more variables: EDU010192D <dbl>, EDU010193D <dbl>, EDU010194D <dbl>,  
#   EDU010195D <dbl>, EDU010196D <dbl>
```

Question 2: Converting to Long Format

Convert the data into long format where each row has only one enrollment value for that Area_name. Display the first 5 rows of your new data set to show that you created this correctly.

```
edu01a_long <- edu01a |>
  pivot_longer(cols = 3:12, #retain area_name + STCOU
               names_to = "EDU_D", #named after unique
               # ending "D" per Q1
               values_to = "Enrollment")

#Displaying first 5 rows below
head(edu01a_long, 5)
```

```
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 EDU010187D    40024299
2 UNITED STATES 00000 EDU010188D    39967624
3 UNITED STATES 00000 EDU010189D    40317775
4 UNITED STATES 00000 EDU010190D    40737600
5 UNITED STATES 00000 EDU010191D    41385442
```

Question 3: Assign Year and State

One of the new columns should now correspond to the old column names that end with a “D”. All columns in these census data files will have this similar format. The first three characters represent the survey with the next four representing the type of value you have from that survey. The last two digits prior to the “D” represent the year of the measurement. For more about the variables see the data information sheet Mastdata.xls).

- Parse the string to pull out the year and convert the year into a numeric value such as 1997 or 2002.
- Grab the first three characters and following four digits to create a new variable representing which measurement was grabbed.
- Hint: Check out the substr() function from base r

```
# Parse the string to pull out year
# It looks like every year is pre-2000, but let's plan for up to 2025
# This assumes there is no data from 1925 or earlier
# Treating year as numeric for now
long_updated <- edu01a_long |> mutate(year = as.numeric(substr(EDU_D, 8,9)),
  measurement = substr(EDU_D, 1,7)) |>
  mutate(year = ifelse(year < 26, year + 2000, year + 1900))

head(long_updated)
```

```
# A tibble: 6 x 6
  area_name      STCOU EDU_D      Enrollment  year measurement
  <chr>          <chr> <chr>          <dbl> <dbl> <chr>
1 UNITED STATES 00000 EDU010187D    40024299  1987 EDU0101
2 UNITED STATES 00000 EDU010188D    39967624  1988 EDU0101
3 UNITED STATES 00000 EDU010189D    40317775  1989 EDU0101
4 UNITED STATES 00000 EDU010190D    40737600  1990 EDU0101
5 UNITED STATES 00000 EDU010191D    41385442  1991 EDU0101
6 UNITED STATES 00000 EDU010192D    42088151  1992 EDU0101
```

Question 4: Split County and Non-County

Create two data sets

- one data set that contains only non-county data
- one data set that contains only county level data

Note that all county measurements have the format “County Name, DD” where “DD” represents the state. This can be used to subset the data. For the county level data, add a class to the tibble called county. Similarly, add a class to the non-county data called state.

```
#For county tibble
county_match <- grep(pattern = ", \\w\\w", long_updated$area_name)
county_tibble <- long_updated[county_match,]
class(county_tibble) <- c("county", class(county_tibble))

#For state tibble
state_match <- grep(pattern = ", \\w\\w", long_updated$area_name, invert = T)
state_tibble <- long_updated[state_match,]
class(state_tibble) <- c("state", class(state_tibble))
```

Print the first 10 rows of each tibble by including `county_tibble` and `state_tibble` in your code chunk.

```
head(county_tibble, 10)
```

```
# A tibble: 10 x 6
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>
1	Autauga, AL	01001	EDU010187D	6829	1987	EDU0101
2	Autauga, AL	01001	EDU010188D	6900	1988	EDU0101
3	Autauga, AL	01001	EDU010189D	6920	1989	EDU0101
4	Autauga, AL	01001	EDU010190D	6847	1990	EDU0101
5	Autauga, AL	01001	EDU010191D	7008	1991	EDU0101
6	Autauga, AL	01001	EDU010192D	7137	1992	EDU0101
7	Autauga, AL	01001	EDU010193D	7152	1993	EDU0101
8	Autauga, AL	01001	EDU010194D	7381	1994	EDU0101
9	Autauga, AL	01001	EDU010195D	7568	1995	EDU0101
10	Autauga, AL	01001	EDU010196D	7834	1996	EDU0101

```
head(state_tibble, 10)
```

```
# A tibble: 10 x 6
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>
1	UNITED STATES	00000	EDU010187D	40024299	1987	EDU0101
2	UNITED STATES	00000	EDU010188D	39967624	1988	EDU0101
3	UNITED STATES	00000	EDU010189D	40317775	1989	EDU0101
4	UNITED STATES	00000	EDU010190D	40737600	1990	EDU0101
5	UNITED STATES	00000	EDU010191D	41385442	1991	EDU0101
6	UNITED STATES	00000	EDU010192D	42088151	1992	EDU0101
7	UNITED STATES	00000	EDU010193D	42724710	1993	EDU0101
8	UNITED STATES	00000	EDU010194D	43369917	1994	EDU0101
9	UNITED STATES	00000	EDU010195D	43993459	1995	EDU0101
10	UNITED STATES	00000	EDU010196D	44715737	1996	EDU0101

Question 5: Assign State to County Tibble

For the county level tibble, create a new variable that describes which state one of these county measurements corresponds to (the two digit abbreviation is fine, see `substr()`).

```
#I prefer to split the string based on delimiter (comma) instead of indexing
#Example
string <- "Autauga, AL"
split <- str_split(string, ",", simplify = TRUE)[-1] # We return a chr matrix,
# and we only care about the last (second) entry

print(split)
```

```
[1] " AL"
```

```
print("Removing space")
```

```
[1] "Removing space"
```

```
clean_split <- str_trim(split)
print(clean_split)
```

```
[1] "AL"
```

```
#Create state variable for county tibble
```

```
county_tibble <- county_tibble |>
  mutate(state = str_trim(str_split(area_name, ",", simplify = TRUE)[-1]))

county_tibble #to show that the addition of the variable was successful
```

```
# A tibble: 31,450 x 7
  area_name STCOU EDU_D Enrollment year measurement state
  <chr>      <chr> <chr>      <dbl> <dbl> <chr>      <chr>
1 Autauga, AL 01001 EDU010187D    6829  1987 EDU0101    AL
2 Autauga, AL 01001 EDU010188D    6900  1988 EDU0101    AL
3 Autauga, AL 01001 EDU010189D    6920  1989 EDU0101    AL
4 Autauga, AL 01001 EDU010190D    6847  1990 EDU0101    AL
5 Autauga, AL 01001 EDU010191D    7008  1991 EDU0101    AL
6 Autauga, AL 01001 EDU010192D    7137  1992 EDU0101    AL
7 Autauga, AL 01001 EDU010193D    7152  1993 EDU0101    AL
8 Autauga, AL 01001 EDU010194D    7381  1994 EDU0101    AL
9 Autauga, AL 01001 EDU010195D    7568  1995 EDU0101    AL
10 Autauga, AL 01001 EDU010196D    7834  1996 EDU0101    AL
# i 31,440 more rows
```

Question 6: Assign Division to State Tibble

For the non-county level tibble, create a new variable called “division” corresponding to the state’s classification of division [here](#). If row corresponds to a non-state (i.e. UNITED STATES), return ERROR for the division. Hint: Use %in% and consider if_else or case_when logic.

Instead of writing ifelse statements manually for every division, we are going to instead read the divisions straight from Wikipedia and assign the correct division to any given state.

We can scrape a Wikipedia table using the rvest package.

Source: [StackOverflow](#)

```
library(rvest) # rvest is in the tidyverse package

# Since we don't want to always have to connect to the url to read our data,
# let's check if we have already saved it
if (file.exists("data/divisions.csv")){
  print("Division data already downloaded from Wikipedia")
  print("Reading .csv file")
  divisions <- read_csv("data/divisions.csv", show_col_types = FALSE)
} else {
  print("No division data found. Downloading from Wikipedia...")
  wiki <- read_html(x =
"https://en.wikipedia.org/wiki/List_of_regions_of_the_United_States",
package="xml2")
  wiki |> html_elements(".wikitable") |> html_table() -> wiki_tables
  # There is only one table, so the first one will give us what we want
  divisions <- wiki_tables[1]
  # Write the file to csv
  write_csv(divisions, file= "data/divisions.csv")
  print("data/divisions.csv successfully created!")
}
```

```
[1] "Division data already downloaded from Wikipedia"
```

```
[1] "Reading .csv file"
```

New names:

```
* `` -> `...1`
```

```
divisions
```

```
# A tibble: 9 x 4
  ...1 Region Division States
<dbl> <chr> <chr> <chr>
1 1 Northeast New England Connecticut Maine Massachusetts New Hampsh~
2 2 Northeast Mid-Atlantic New Jersey New York Pennsylvania
3 3 Midwest East North Central Illinois Indiana Michigan Ohio Wisconsin
4 4 Midwest West North Central Iowa Kansas Minnesota Missouri Nebraska No~
5 5 South South Atlantic Delaware District of Columbia Florida Geor~
6 6 South East South Central Alabama Kentucky Mississippi Tennessee
7 7 South West South Central Arkansas Louisiana Oklahoma Texas
8 8 West Mountain Arizona Colorado Idaho Montana Nevada New ~
9 9 West Pacific Alaska California Hawaii Oregon Washington
```

Note how all states in any given region are stored in the same cell, separated by spaces. We can either transform the States column by splitting up the states or leave as is and process the state correctly when reading our other datasets.

We can filter columns by the state in our divisions tibble by using `if_any` and `str_detect`.

```
# Make uppercase to improve matching
divisions$States <-divisions$States |> toupper()

get_division_for_state <- function(state_name){

  # Check for the state name in the divisions df and filter
  # Assumes state only appears once in the tibble
  # Add word boundaries to our regex to avoid substring matching
  # E.g "Kansas" shouldn't match "Arkansas"
  match_pattern <- paste0("\\b", toupper(state_name), "\\b")
  division_row <- divisions |>
    filter(if_any(States, ~str_detect(.x, match_pattern)))
  division <- division_row$Division
  # Return "ERROR" if there is no match to state
  if (length(division) == 0){
    return ("ERROR")
  }
  else {
    return (division)
  }
}
```



```
# Apply our function to the non county tibble
```

```
state_tibble_test <- state_tibble |> mutate(division =  
  map_chr(area_name, get_division_for_state))  
tail(state_tibble_test)
```

```
# A tibble: 6 x 7
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement	division
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>
1	WYOMING	56000	EDU010191D	98782	1991	EDU0101	Mountain
2	WYOMING	56000	EDU010192D	101715	1992	EDU0101	Mountain
3	WYOMING	56000	EDU010193D	100729	1993	EDU0101	Mountain
4	WYOMING	56000	EDU010194D	100899	1994	EDU0101	Mountain
5	WYOMING	56000	EDU010195D	100369	1995	EDU0101	Mountain
6	WYOMING	56000	EDU010196D	99859	1996	EDU0101	Mountain

Function Wrapping

Function 1: Step 1 & Step 2

Write one function that combines Steps 1 and 2 above. Give an optional argument (that is it has a default value) that allows the user to specify the name of the column representing the value (enrollment for these data sets).

```
select_and_convert <- function(data_path_in_quotes, value_colname = "Enrollment"){  
  edu <- read_csv(data_path_in_quotes,  
    col_select = c(Area_name, STCOU, ends_with("D")),  
    show_col_types = FALSE) |>  
  rename(  
    area_name = Area_name  
  )  
  
  edu_long <- edu |>  
    pivot_longer(cols = 3:12,  
      names_to = "EDU_D",  
      values_to = value_colname)  
  print(head(edu_long, 5))  
  return(edu_long)  
}
```

```
#making sure the function works
function1 <- select_and_convert("data/EDU01b.csv")
```

```
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 EDU010197D    44534459
2 UNITED STATES 00000 EDU010198D    46245814
3 UNITED STATES 00000 EDU010199D    46368903
4 UNITED STATES 00000 EDU010200D    46818690
5 UNITED STATES 00000 EDU010201D    47127066
```

Function 2: Step 3

Write a function that takes the output from Step 2 and performs Step 3

```
get_year_and_measurement <-function(long_data){
  print("Updating long data with year and measurement")
  long_data_updated <- long_data |>
    mutate(year = as.numeric(substr(EDU_D, 8,9)),
           measurement = substr(EDU_D, 1,7)) |>
    mutate(year = ifelse(year < 26, year + 2000, year + 1900))

  return (long_data_updated)
}

get_year_and_measurement(function1) #to make sure the function works
```

```
[1] "Updating long data with year and measurement"
```

```
# A tibble: 31,980 x 6
  area_name      STCOU EDU_D      Enrollment  year measurement
  <chr>          <chr> <chr>          <dbl> <dbl> <chr>
1 UNITED STATES 00000 EDU010197D    44534459  1997 EDU0101
2 UNITED STATES 00000 EDU010198D    46245814  1998 EDU0101
3 UNITED STATES 00000 EDU010199D    46368903  1999 EDU0101
4 UNITED STATES 00000 EDU010200D    46818690  2000 EDU0102
5 UNITED STATES 00000 EDU010201D    47127066  2001 EDU0102
6 UNITED STATES 00000 EDU010202D    47606570  2002 EDU0102
7 UNITED STATES 00000 EDU015203D    48506317  2003 EDU0152
```

```

8 UNITED STATES 00000 EDU015204D 48693287 2004 EDU0152
9 UNITED STATES 00000 EDU015205D 48978555 2005 EDU0152
10 UNITED STATES 00000 EDU015206D 49140702 2006 EDU0152
# i 31,970 more rows

```

Function 3: Step 5

Write a function to do Step 5

```

get_state <- function(county_tibble){
  print("Assigning State to county tibble")
  county_tibble_with_state <- county_tibble |>
  mutate(state = str_trim(str_split(area_name, ",", simplify = TRUE)[,-1]))

  return(county_tibble_with_state)
}

```

Function 4: Step 6

Write a function to do step 6

```

get_division <- function(state_tibble){
  print("Assigning division to state tibble")
  state_tibble_with_division <- state_tibble |>
  mutate(division = map_chr(area_name, get_division_for_state))

  return(state_tibble_with_division)
}

```

Function 5: Step 4

Write another function that takes in the output from Step 3 and creates the two tibbles in Step 4, calls the above two functions (to perform Steps 5 and 6), and returns two final tibbles.

```

returning_final_tibbles <- function(long_data_updated){
  county_match <- grep(pattern = ", \\w\\w", long_data_updated$area_name)
  county_tibble <- long_data_updated[county_match,]
  class(county_tibble) <- c("county", class(county_tibble))

  state_match <- grep(pattern = ", \\w\\w", long_data_updated$area_name, invert = T)

```

```

state_tibble <- long_data_updated[state_match,]
class(state_tibble) <- c("state", class(state_tibble))

county_tibble_final <- get_state(county_tibble)
state_tibble_final <- get_division(state_tibble)

return(list(county_tibble_final, state_tibble_final))
}

#making sure the function works
returning_final_tibbles(get_year_and_measurement(function1))

```

```

[1] "Updating long data with year and measurement"
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

```
[[1]]
```

```
# A tibble: 31,450 x 7
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement	state
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>
1	Autauga, AL	01001	EDU010197D	8099	1997	EDU0101	AL
2	Autauga, AL	01001	EDU010198D	8211	1998	EDU0101	AL
3	Autauga, AL	01001	EDU010199D	8489	1999	EDU0101	AL
4	Autauga, AL	01001	EDU010200D	8912	2000	EDU0102	AL
5	Autauga, AL	01001	EDU010201D	8626	2001	EDU0102	AL
6	Autauga, AL	01001	EDU010202D	8762	2002	EDU0102	AL
7	Autauga, AL	01001	EDU015203D	9105	2003	EDU0152	AL
8	Autauga, AL	01001	EDU015204D	9200	2004	EDU0152	AL
9	Autauga, AL	01001	EDU015205D	9559	2005	EDU0152	AL
10	Autauga, AL	01001	EDU015206D	9652	2006	EDU0152	AL

```
# i 31,440 more rows
```

```
[[2]]
```

```
# A tibble: 530 x 7
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement	division
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>
1	UNITED STATES	00000	EDU010197D	44534459	1997	EDU0101	ERROR
2	UNITED STATES	00000	EDU010198D	46245814	1998	EDU0101	ERROR
3	UNITED STATES	00000	EDU010199D	46368903	1999	EDU0101	ERROR
4	UNITED STATES	00000	EDU010200D	46818690	2000	EDU0102	ERROR
5	UNITED STATES	00000	EDU010201D	47127066	2001	EDU0102	ERROR

```

6 UNITED STATES 00000 EDU010202D 47606570 2002 EDU0102 ERROR
7 UNITED STATES 00000 EDU015203D 48506317 2003 EDU0152 ERROR
8 UNITED STATES 00000 EDU015204D 48693287 2004 EDU0152 ERROR
9 UNITED STATES 00000 EDU015205D 48978555 2005 EDU0152 ERROR
10 UNITED STATES 00000 EDU015206D 49140702 2006 EDU0152 ERROR
# i 520 more rows

```

Wrap Everything in One Function Call (Wrapper Function)

```

clean_data_wrapper <- function(url, value = "Enrollment"){
  result <- select_and_convert(url, value_colname = value) |>
    get_year_and_measurement() |>
    returning_final_tibbles()
}

```

Call It and Combine Data

Call the function you made two times to read in and parse the two .csv files mentioned so far. Be sure to call the new value column the same in both function calls.

```

data_a <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/EDU01a.csv")

```

```

[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 EDU010187D    40024299
2 UNITED STATES 00000 EDU010188D    39967624
3 UNITED STATES 00000 EDU010189D    40317775
4 UNITED STATES 00000 EDU010190D    40737600
5 UNITED STATES 00000 EDU010191D    41385442
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

```

data_b <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/EDU01b.csv")

```

```

[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name STCOU EDU_D Enrollment
  <chr>      <chr> <chr>      <dbl>
1 UNITED STATES 00000 EDU010197D 44534459
2 UNITED STATES 00000 EDU010198D 46245814
3 UNITED STATES 00000 EDU010199D 46368903
4 UNITED STATES 00000 EDU010200D 46818690
5 UNITED STATES 00000 EDU010201D 47127066
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

Write a single short function that takes in the results of two calls to your wrapper function. The function should combine the tibbles appropriately (that is the two county level data sets get combined and the two non-county level data sets get combined). This can easily be done within your function using some calls to `dplyr::bind_rows()`. The function should then return two data sets as one object (in the same format as the input data sets as we will be combining this output with more calls to the wrapper function in a bit).

```

combining_tibbles <- function(tibble1, tibble2){
  county_tibbles_combined <- bind_rows(tibble1[[1]], tibble2[[1]])
  state_tibbles_combined <- bind_rows(tibble1[[2]], tibble2[[2]])
  return(list(county_tibbles_combined, state_tibbles_combined))
}

```

Call this function to combine the result of the two calls to the wrapper function.

```

#saving this test to use it for testing plots later
test_data <- combining_tibbles(data_a, data_b)

test_data #to show that it worked

```

```

[[1]]
# A tibble: 62,900 x 7
  area_name STCOU EDU_D Enrollment year measurement state
  <chr>      <chr> <chr>      <dbl> <dbl> <chr>      <chr>
1 Autauga, AL 01001 EDU010187D 6829 1987 EDU0101 AL
2 Autauga, AL 01001 EDU010188D 6900 1988 EDU0101 AL
3 Autauga, AL 01001 EDU010189D 6920 1989 EDU0101 AL
4 Autauga, AL 01001 EDU010190D 6847 1990 EDU0101 AL
5 Autauga, AL 01001 EDU010191D 7008 1991 EDU0101 AL

```

```

6 Autauga, AL 01001 EDU010192D      7137  1992 EDU0101      AL
7 Autauga, AL 01001 EDU010193D      7152  1993 EDU0101      AL
8 Autauga, AL 01001 EDU010194D      7381  1994 EDU0101      AL
9 Autauga, AL 01001 EDU010195D      7568  1995 EDU0101      AL
10 Autauga, AL 01001 EDU010196D     7834  1996 EDU0101      AL
# i 62,890 more rows

```

```
[[2]]
```

```
# A tibble: 1,060 x 7
```

	area_name	STCOU	EDU_D	Enrollment	year	measurement	division
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>
1	UNITED STATES	00000	EDU010187D	40024299	1987	EDU0101	ERROR
2	UNITED STATES	00000	EDU010188D	39967624	1988	EDU0101	ERROR
3	UNITED STATES	00000	EDU010189D	40317775	1989	EDU0101	ERROR
4	UNITED STATES	00000	EDU010190D	40737600	1990	EDU0101	ERROR
5	UNITED STATES	00000	EDU010191D	41385442	1991	EDU0101	ERROR
6	UNITED STATES	00000	EDU010192D	42088151	1992	EDU0101	ERROR
7	UNITED STATES	00000	EDU010193D	42724710	1993	EDU0101	ERROR
8	UNITED STATES	00000	EDU010194D	43369917	1994	EDU0101	ERROR
9	UNITED STATES	00000	EDU010195D	43993459	1995	EDU0101	ERROR
10	UNITED STATES	00000	EDU010196D	44715737	1996	EDU0101	ERROR

```

# i 1,050 more rows

```

Write Generic Functions for Summarizing

Plotting State Data

Let's show commas in our y-axis to make it more readable.

Source: [StackOverflow](#)

```

# We will use the library scales
library(scales)

```

```

plot.state <- function(df, var_name = "Enrollment"){
  # Create title base on our supplied var name
  plot_title <- paste0("Mean ", var_name, " by Division")

  df |>
    filter(division != "ERROR") |>
    group_by(division, year) |>
    summarize(mean_enrollment = mean(get(var_name))) |>

```

```

mutate(division = as.factor(division)) |>

# Plotting functions
ggplot(aes(year, mean_enrollment, color = division)) +
  geom_line() +
  labs(title = plot_title, x = "Year", y = paste0("Mean ", var_name)) +
  guides(color = guide_legend("U.S. Division")) + # Rename Legend
  scale_y_continuous(label=comma)
}

```

Plotting County Data

```

plot.county <- function(df, var_name = "Enrollment",
                        state = "NC",
                        top_or_bottom = "top", n = 5){
  # Argument validation
  # Try to match by state
  if (is.na(state.name[match(state, state.abb)])){
    stop("Argument Error: state must be two letter state abb, e.g. 'NC' ")
  }
  if (!(all.equal(n, as.integer(n))) == TRUE ){
    stop("Argument Error: Please use an integer for n")
  }

  # Create title based on our supplied var name
  plot_title <- paste0(ifelse(top_or_bottom == "top", "Highest", "Lowest"), " ",
                        var_name, " in ", state, " by County")

  # Helper function, not sure if this is the most efficient way to handle this
  display_function <- function(df, col, top_or_bottom){
    if (top_or_bottom == "top"){
      df |> arrange(desc({{col}})) # Nested brackets to refer to the column
    }
    else if(top_or_bottom == "bottom"){
      df |> arrange({{col}})
    }
    else
      stop("Argument Error: top_or_bottom must be 'top' or 'bottom'")
  }
}

```



```

}

# Get n top or bottom counties
counties <- df |>
  filter({{state}} == state ) |>
  group_by(area_name) |>
  summarize(mean_enrollment = mean(get(var_name))) |>
  display_function(mean_enrollment, top_or_bottom) |>
  head(n)

# Filter df by counties and plot
# 2 cols for legend so that it's not cut off at the top
df |> filter(df$area_name %in% counties$area_name) |> group_by(year, area_name) |>
  ggplot(aes(year, get(var_name), color = area_name)) + geom_line() +
  labs(title = plot_title, x = "Year", y = paste0(var_name)) +
  guides(color = guide_legend("Location", ncol = 2)) +
  scale_y_continuous(label=comma)

}

```

Put It Together

Run your data processing function on the two enrollment URLs given previously, specifying an appropriate name for the enrollment data column.

```
data1 <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/EDU01a.csv")
```

```

[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name STCOU EDU_D Enrollment
  <chr>      <chr> <chr>      <dbl>
1 UNITED STATES 00000 EDU010187D 40024299
2 UNITED STATES 00000 EDU010188D 39967624
3 UNITED STATES 00000 EDU010189D 40317775
4 UNITED STATES 00000 EDU010190D 40737600
5 UNITED STATES 00000 EDU010191D 41385442
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

```
data2 <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/EDU01b.csv")
```

```
[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>         <chr> <chr>         <dbl>
1 UNITED STATES 00000 EDU010197D    44534459
2 UNITED STATES 00000 EDU010198D    46245814
3 UNITED STATES 00000 EDU010199D    46368903
4 UNITED STATES 00000 EDU010200D    46818690
5 UNITED STATES 00000 EDU010201D    47127066
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"
```

Run your data combining function to put these into one object (with two data frames)

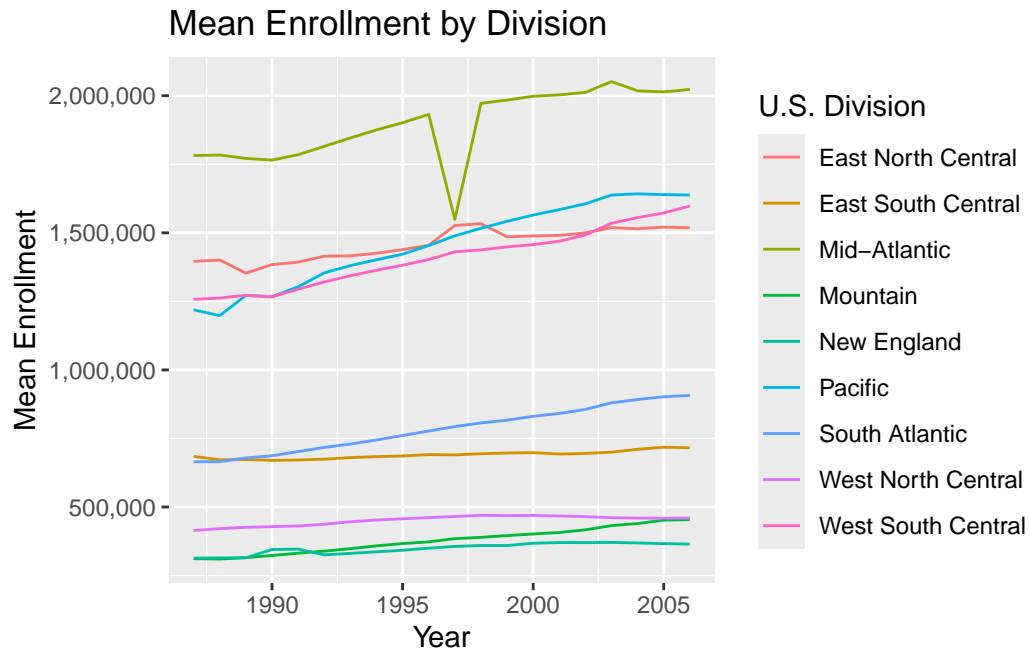
```
one_object <- combining_tibbles(data1, data2)
```

(Use appropriate indexing (ex. `[[1]]`) to reference the correct data frame)

Use the plot function on the state data frame

```
plot(one_object[[2]])
```

``summarise()`` has grouped output by 'division'. You can override using the ``groups`` argument.

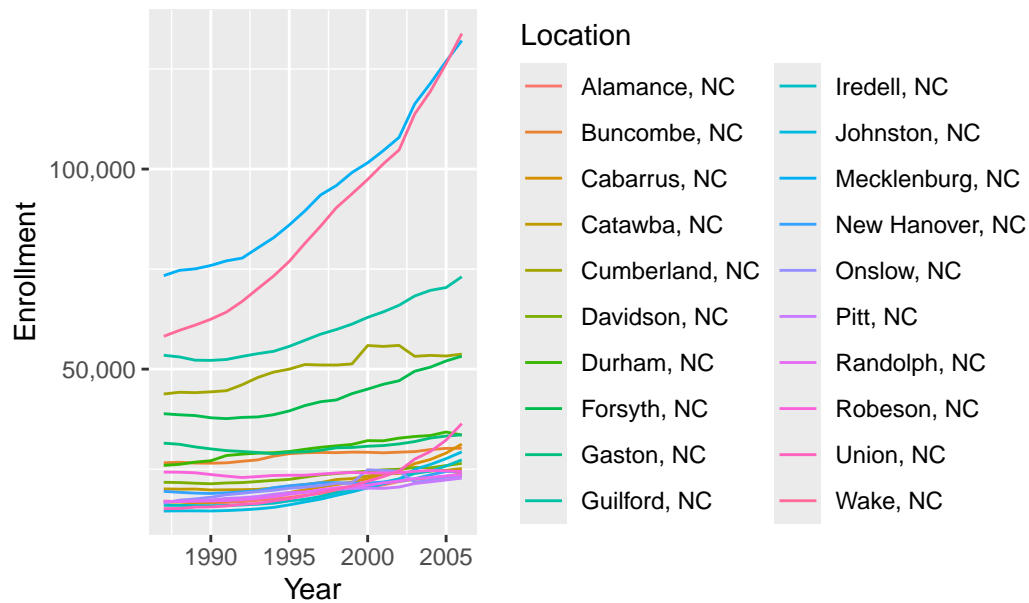


Use the plot function on the county data frame

- Once specifying the state to be “NC”, the group being the top, the number looked at being 20

```
plot(one_object[[1]], top_or_bottom = "top", n = 20, state = "NC")
```

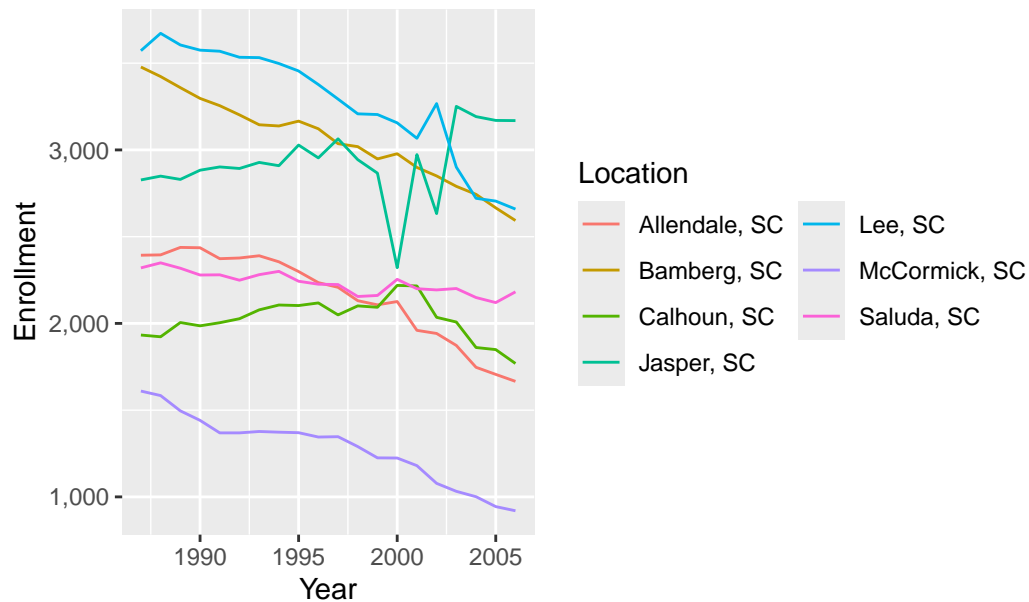
Highest Enrollment in NC by County



– Once specifying the state to be “SC”, the group being the bottom, the number looked at being 7

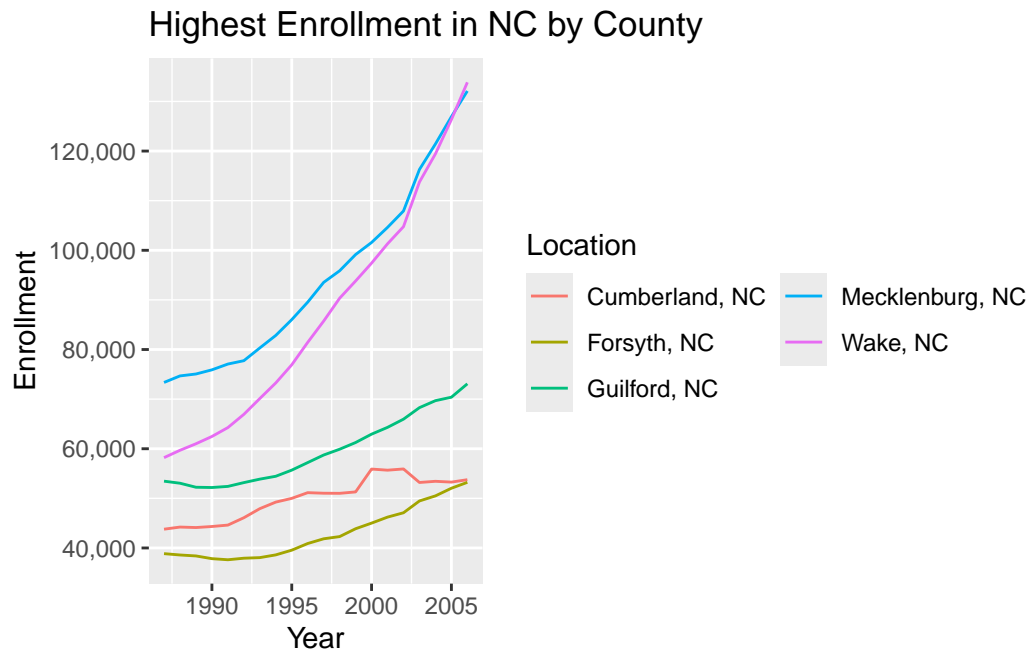
```
plot(one_object[[1]], top_or_bottom = "bottom", n = 7, state = "SC")
```

Lowest Enrollment in SC by County



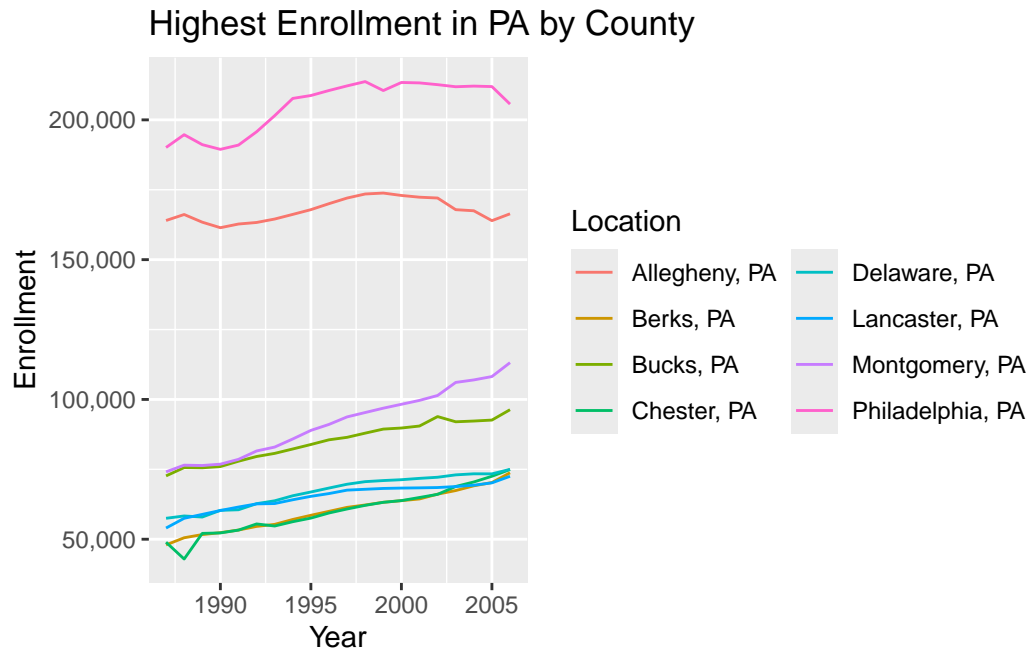
- Once without specifying anything (defaults used)

```
plot(one_object[[1]])
```



- Once specifying the state to be “PA”, the group being the top, the number looked at being 8

```
plot(one_object[[1]], top_or_bottom = "top", n = 8, state = "PA")
```



Lastly, read in another couple similar data sets and apply your functions!

Run your data processing function on the four data sets at URLs given.

```
dataPa <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/PST01a.csv")
```

```
[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name STCOU EDU_D Enrollment
  <chr>      <chr> <chr>      <dbl>
1 UNITED STATES 00000 PST015171D 206827028
2 UNITED STATES 00000 PST015172D 209283904
3 UNITED STATES 00000 PST015173D 211357490
4 UNITED STATES 00000 PST015174D 213341552
5 UNITED STATES 00000 PST015175D 215465246
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"
```

```
dataPb <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/PST01b.csv")
```

```
[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
```

```

  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 PST025182D 231665106
2 UNITED STATES 00000 PST025183D 233792697
3 UNITED STATES 00000 PST025184D 235825544
4 UNITED STATES 00000 PST025185D 237924311
5 UNITED STATES 00000 PST025186D 240133472
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

```
dataPc <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/PST01c.csv")
```

```

[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 PST035191D 252980941
2 UNITED STATES 00000 PST035192D 256514224
3 UNITED STATES 00000 PST035193D 259918588
4 UNITED STATES 00000 PST035194D 263125821
5 UNITED STATES 00000 PST035195D 266278393
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

```
dataPd <- clean_data_wrapper("https://www4.stat.ncsu.edu/~online/datasets/PST01d.csv")
```

```

[1] "Updating long data with year and measurement"
# A tibble: 5 x 4
  area_name      STCOU EDU_D      Enrollment
  <chr>          <chr> <chr>          <dbl>
1 UNITED STATES 00000 PST045200D 282171957
2 UNITED STATES 00000 PST045201D 285081556
3 UNITED STATES 00000 PST045202D 287803914
4 UNITED STATES 00000 PST045203D 290326418
5 UNITED STATES 00000 PST045204D 293045739
[1] "Assigning State to county tibble"
[1] "Assigning division to state tibble"

```

Run your data combining function (probably three times) to put these into one object (with two data frames)

```
once <- combining_tibbles(dataPa, dataPb)


twice <- combining_tibbles(once, dataPc)

thrice <- combining_tibbles(twice, dataPd)
```

Use the plot function on the state data frame


```
plot(thrice[[2]])
```

`summarise()` has grouped output by 'division'. You can override using the `groups` argument.




Use the plot function on the county data frame – Once specifying the state to be “CA”, the group being the top, the number looked at being 15

```
plot(thrice[[1]], top_or_bottom = "top", n = 15, state = "CA")
```


– Once specifying the state to be “TX”, the group being the top, the number looked at being 4

```
plot(thrice[[1]], top_or_bottom = "top", n = 4, state = "TX")
```



- Once without specifying anything (defaults used)

```
plot(thrice[[1]])
```



- Once specifying the state to be “NY”, the group being the top, the number looked at being 10

```
plot(thrice[[1]], top_or_bottom = "top", n = 10, state = "NY")
```

Highest Enrollment in NY by County

